


Homework – 5

Question 12.1

- (a) Given the classification imbalance in hepatic injury status, describe how you would create a training and testing set.

Answer: On observing the hepatic dataset, we observe that there is a huge imbalance in injury status. In order to overcome this, we can use stratified random sampling approach. Below graph shows the imbalance between the three injury classes.



- (b) Which classification statistic would you choose to optimize for this exercise and why?

Answer: For this exercise, we will be using “Accuracy” as the classification statistic as Accuracy is the percentage of correctly classified instances out of all instances.

- (c) Split the data into a training and a testing set, pre-process the data, and build models described in this chapter for the biological predictors. Using each model to predict on the testing set, which model has the best predictive ability for the biological predictors and what is the optimal performance?

Answer: We split our data into 75% training and 25% as testing data. Before splitting the data, we preprocess it by removing near-zero variance predictors, center and scale, and by removing highly correlated predictors.

Pre-Process:

```

#PreProcess the data

#Removing nearzero variance
NV <- nearZeroVar(bio)
NV
noZVbio <- bio[,-NV]
noZVbio

#Calculating missing values
sum(is.na(noZVbio))

#Finding correlation
set.seed(1)
highCorBio<-findCorrelation(cor(noZVbio),cutoff = .80)
filteredCorBio <- noZVbio[,-highCorBio]

```

Splitting Data:

```

# Split the data

set.seed(1)

trainingRows = createDataPartition(injury, p = .75, list= FALSE)

trainBio <- filteredCorBio[ trainingRows, ]
testBio <- filteredCorBio[-trainingRows, ]

trainInjury <- injury[trainingRows]
testInjury <- injury[-trainingRows]

```

Logistic Regression Model:

Confusion Matrix and Statistics

	Reference		
Prediction	Mild	None	Severe
Mild	21	9	3
None	14	11	1
Severe	1	6	3

Overall Statistics

```

Accuracy : 0.5072
95% CI : (0.3841, 0.6298)
No Information Rate : 0.5217
P-Value [Acc > NIR] : 0.6416

```

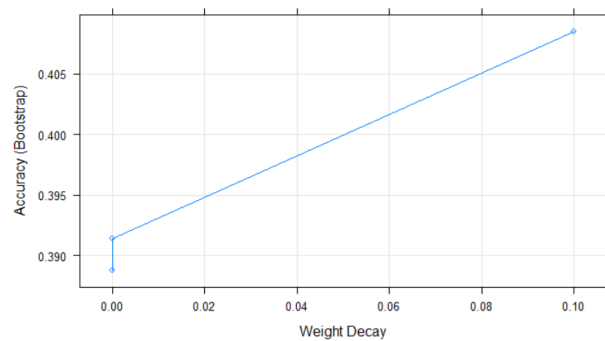
Kappa : 0.1701

Mcnemar's Test P-Value : 0.1295

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.5833	0.4231	0.42857
Specificity	0.6364	0.6512	0.88710
Pos Pred Value	0.6364	0.4231	0.30000
Neg Pred Value	0.5833	0.6512	0.93220
Prevalence	0.5217	0.3768	0.10145
Detection Rate	0.3043	0.1594	0.04348
Detection Prevalence	0.4783	0.3768	0.14493
Balanced Accuracy	0.6098	0.5371	0.65783

Plot for Logistic Regression:



Linear Discriminant Analysis Model:

Confusion Matrix and Statistics

		Reference		
Prediction		Mild	None	Severe
Mild		20	9	6
None		14	14	0
Severe		2	3	1

Overall Statistics

Accuracy : 0.5072
95% CI : (0.3841, 0.6298)
No Information Rate : 0.5217
P-Value [Acc > NIR] : 0.6416

Kappa : 0.141

McNemar's Test P-Value : 0.1075

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.5556	0.5385	0.14286
Specificity	0.5455	0.6744	0.91935
Pos Pred Value	0.5714	0.5000	0.16667
Neg Pred Value	0.5294	0.7073	0.90476
Prevalence	0.5217	0.3768	0.10145
Detection Rate	0.2899	0.2029	0.01449
Detection Prevalence	0.5072	0.4058	0.08696
Balanced Accuracy	0.5505	0.6064	0.53111

Partial Least Square Discriminant Analysis Model:

Confusion Matrix and Statistics

		Reference		
Prediction		Mild	None	Severe
Mild		27	16	7
None		9	10	0
Severe		0	0	0

Overall Statistics

Accuracy : 0.5362
95% CI : (0.412, 0.6572)
No Information Rate : 0.5217
P-Value [Acc > NIR] : 0.4528

Kappa : 0.105

McNemar's Test P-Value : NA

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.7500	0.3846	0.0000
Specificity	0.3030	0.7907	1.0000
Pos Pred Value	0.5400	0.5263	NaN
Neg Pred Value	0.5263	0.6800	0.8986
Prevalence	0.5217	0.3768	0.1014
Detection Rate	0.3913	0.1449	0.0000
Detection Prevalence	0.7246	0.2754	0.0000
Balanced Accuracy	0.5265	0.5877	0.5000

Penalized Model:

Confusion Matrix and Statistics

		Reference		
Prediction		Mild	None	Severe
Mild		30	17	7
None		6	9	0
Severe		0	0	0

Overall Statistics

Accuracy : 0.5652
95% CI : (0.4404, 0.6842)
No Information Rate : 0.5217
P-Value [Acc > NIR] : 0.274

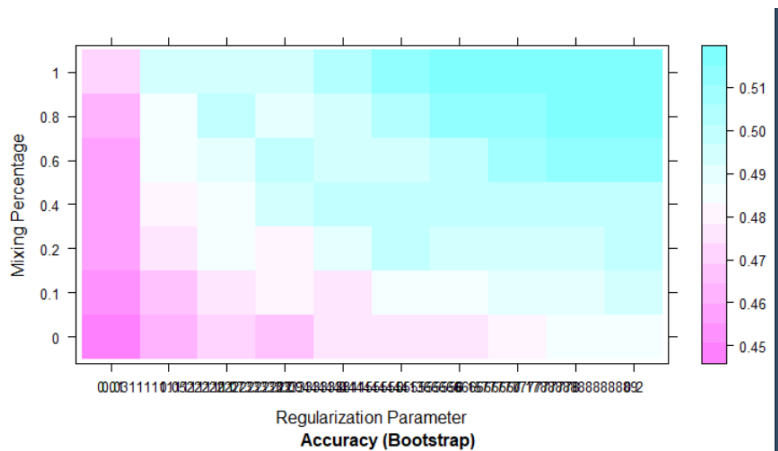
Kappa : 0.1471

McNemar's Test P-Value : NA

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.8333	0.3462	0.0000
Specificity	0.2727	0.8605	1.0000
Pos Pred Value	0.5556	0.6000	NaN
Neg Pred Value	0.6000	0.6852	0.8986
Prevalence	0.5217	0.3768	0.1014
Detection Rate	0.4348	0.1304	0.0000
Detection Prevalence	0.7826	0.2174	0.0000
Balanced Accuracy	0.5530	0.6033	0.5000

Plot of Penalized Model:



Nearest Shrunken Centroids:

Confusion Matrix and Statistics

	Reference		
Prediction	Mild	None	Severe
Mild	36	26	7
None	0	0	0
Severe	0	0	0

Overall Statistics

Accuracy : 0.5217
95% CI : (0.398, 0.6435)
No Information Rate : 0.5217
P-Value [Acc > NIR] : 0.5486

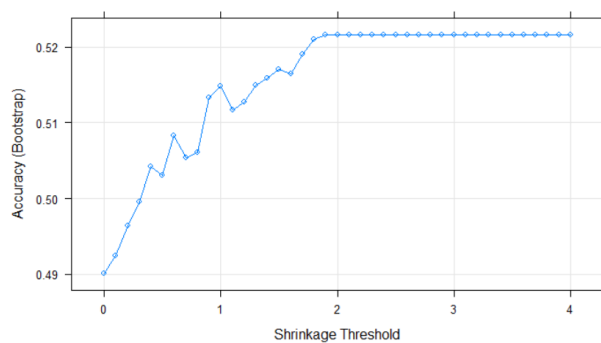
Kappa : 0

McNemar's Test P-Value : NA

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	1.0000	0.0000	0.0000
Specificity	0.0000	1.0000	1.0000
Pos Pred Value	0.5217	NaN	NaN
Neg Pred Value	NaN	0.6232	0.8986
Prevalence	0.5217	0.3768	0.1014
Detection Rate	0.5217	0.0000	0.0000
Detection Prevalence	1.0000	0.0000	0.0000
Balanced Accuracy	0.5000	0.5000	0.5000

Plot of Nearest Shrunken Centroids:




Model	Accuracy
Logistic Regression	0.5072
Linear Discriminant Analysis	0.5072
Partial Least Square Discriminant Analysis	0.5362
Penalized Model	0.5652
Nearest Shrunken Centroids	0.5217

Out of all the models **Penalized Model** has the maximum accuracy of **0.5652**.

(d) For the optimal model for the biological predictors, what are the top five important predictors?

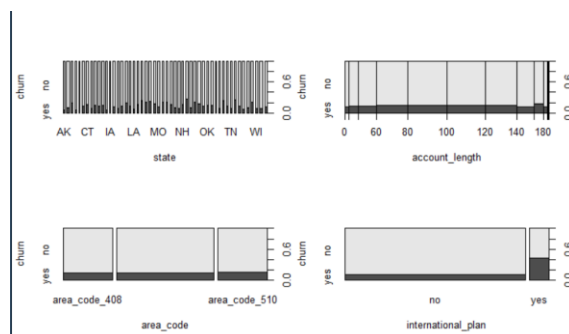
Answer: The top five important predictors of Penalized Model are:

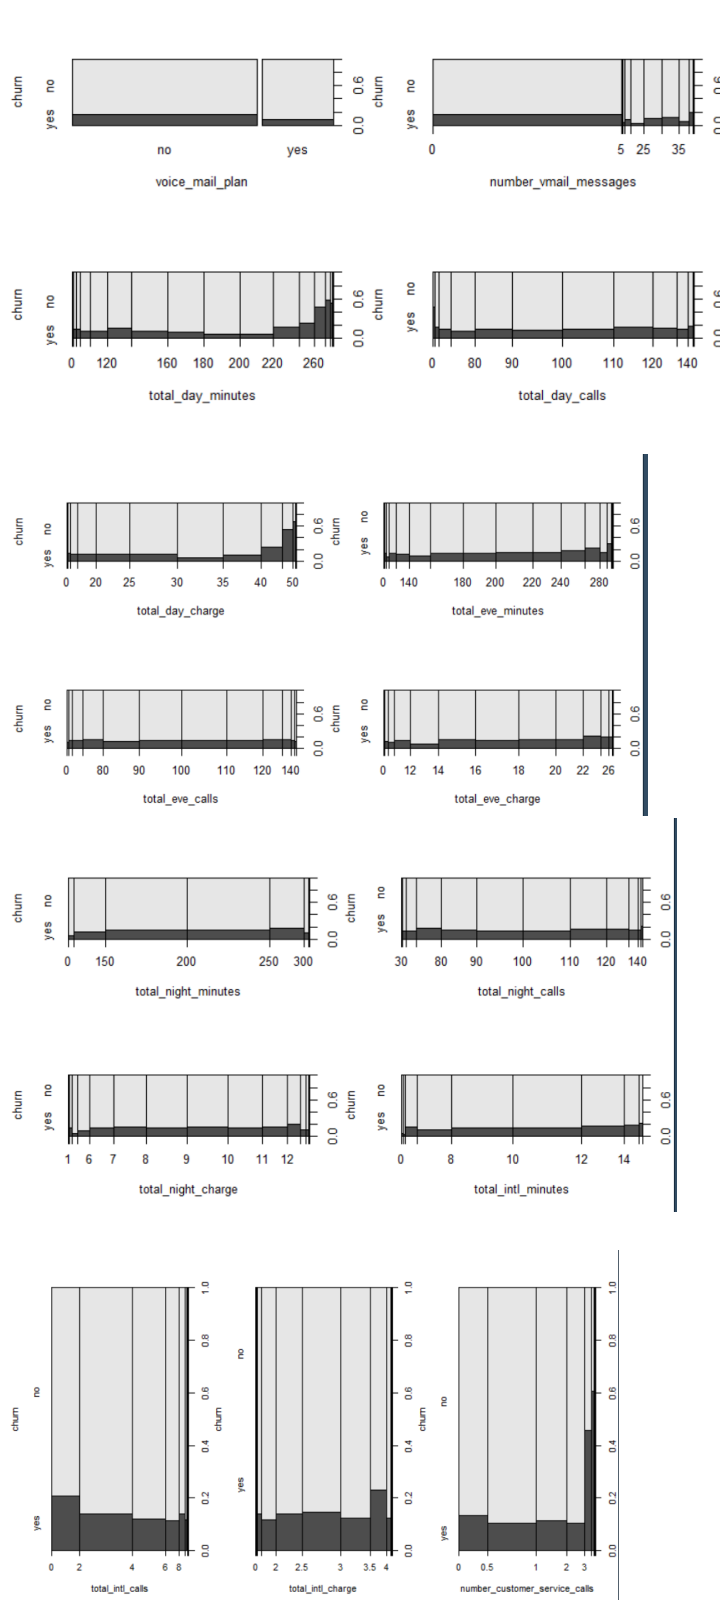


Question 12.3

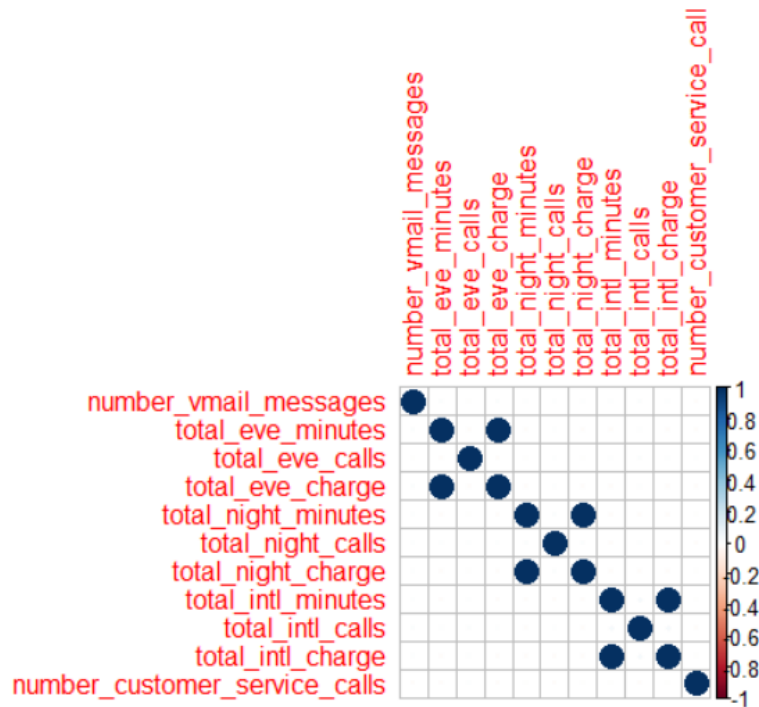
- (a) Explore the data by visualizing the relationship between the predictors and the outcome. Are there important features of the predictor data themselves, such as between-predictor correlations or degenerate distributions?

Answer: We plot each predictor against the outcome to understand the relationship between them:





We also find the correlation between the predictors using the corplot.



(b) What criteria should be used to evaluate the effectiveness of the models?

Answer: For this question, we will be using “**ROC**” as the classification statistic as we are using **LGOCV** as the resampling technique.

(c) Fit models covered in class to the training set and tune them via resampling. Which model has the best performance?

Answer:

Pre-Process: We remove near zero variance predictors, center and scale the data, correct the skewness using BoxCox and remove highly correlated data.

```
#PreProcess

newdata <- preProcess(predict_train, method = c("center", "scale", "BoxCox"))
newdata_train <- predict(newdata, predict_train)
newdata_test <- predict(newdata, predict_test)

NV3 <- nearZeroVar(newdata_train)
newdata_train <- newdata_train[-NV3]
newdata_test <- newdata_test[-NV3]

highcor <- cor(newdata_train)
highcorpredict <- findCorrelation(highcor)

newdata_train <- newdata_train[,-highcorpredict]
newdata_test <- newdata_test[,-highcorpredict]
```

Logistic Regression Model:

```
Generalized Linear Model
3333 samples
 14 predictor
 2 classes: 'yes', 'no'

Pre-processing: centered (14), scaled (14)
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 2501, 2501, 2501, 2501, 2501, 2501, ...
Resampling results:

ROC      Sens  Spec
0.5178488 0     1
```

Linear Discriminant Analysis Model:

```
> lda
Linear Discriminant Analysis
3333 samples
 14 predictor
 2 classes: 'yes', 'no'

Pre-processing: centered (14), scaled (14)
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 2501, 2501, 2501, 2501, 2501, 2501, ...
Resampling results:

ROC      Sens  Spec
0.5179583 0     1
```

Partial Least Square Discriminant Analysis Model:

```
> plsFitc
Partial Least Squares
3333 samples
 14 predictor
 2 classes: 'yes', 'no'

Pre-processing: centered (14), scaled (14)
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 2501, 2501, 2501, 2501, 2501, 2501, ...
Resampling results:

ROC      Sens  Spec
0.5200023 0     1

Tuning parameter 'ncomp' was held constant at a value of 1
```

Penalized Model:

Resampling results across tuning parameters:

alpha	lambda	ROC	Sens	Spec
0.0	0.01000000	0.5181737	0	1
0.0	0.03111111	0.5185520	0	1
0.0	0.05222222	0.5188450	0	1
0.0	0.07333333	0.5190398	0	1
0.0	0.09444444	0.5191742	0	1
0.0	0.11555556	0.5192664	0	1
0.0	0.13666667	0.5193272	0	1
0.0	0.15777778	0.5194139	0	1
0.0	0.17888889	0.5194448	0	1
0.0	0.20000000	0.5195239	0	1
0.1	0.01000000	0.5178146	0	1
0.1	0.03111111	0.5144537	0	1
0.1	0.05222222	0.5083886	0	1
0.1	0.07333333	0.5025625	0	1
0.1	0.09444444	0.4981894	0	1
0.1	0.11555556	0.4921711	0	1
0.1	0.13666667	0.4934438	0	1
0.1	0.15777778	0.4943153	0	1
0.1	0.17888889	0.4984586	0	1
0.1	0.20000000	0.5000000	0	1
0.2	0.01000000	0.5164733	0	1
0.2	0.03111111	0.5056959	0	1
0.2	0.05222222	0.4960028	0	1
0.2	0.07333333	0.4930243	0	1
0.2	0.09444444	0.4984841	0	1
0.2	0.11555556	0.5000000	0	1
0.2	0.13666667	0.5000000	0	1
0.2	0.15777778	0.5000000	0	1
0.2	0.17888889	0.5000000	0	1
0.2	0.20000000	0.5000000	0	1
0.4	0.01000000	0.5119995	0	1
0.4	0.03111111	0.4926297	0	1
0.4	0.05222222	0.5000000	0	1
0.4	0.07333333	0.5000000	0	1
0.4	0.09444444	0.5000000	0	1
0.4	0.11555556	0.5000000	0	1
0.4	0.13666667	0.5000000	0	1
0.4	0.15777778	0.5000000	0	1
0.4	0.17888889	0.5000000	0	1
0.4	0.20000000	0.5000000	0	1
0.6	0.01000000	0.5061402	0	1
0.6	0.03111111	0.4984841	0	1
0.6	0.05222222	0.5000000	0	1
0.6	0.07333333	0.5000000	0	1
0.6	0.09444444	0.5000000	0	1
0.6	0.11555556	0.5000000	0	1
0.6	0.13666667	0.5000000	0	1
0.6	0.15777778	0.5000000	0	1
0.6	0.17888889	0.5000000	0	1
0.6	0.20000000	0.5000000	0	1
0.8	0.01000000	0.5010257	0	1
0.8	0.03111111	0.5000000	0	1
0.8	0.05222222	0.5000000	0	1
0.8	0.07333333	0.5000000	0	1
0.8	0.09444444	0.5000000	0	1
0.8	0.11555556	0.5000000	0	1
0.8	0.13666667	0.5000000	0	1
0.8	0.15777778	0.5000000	0	1

0.8	0.17888889	0.5000000	0	1
0.8	0.20000000	0.5000000	0	1
1.0	0.01000000	0.4961723	0	1
1.0	0.03111111	0.5000000	0	1
1.0	0.05222222	0.5000000	0	1
1.0	0.07333333	0.5000000	0	1
1.0	0.09444444	0.5000000	0	1
1.0	0.11555556	0.5000000	0	1
1.0	0.13666667	0.5000000	0	1
1.0	0.15777778	0.5000000	0	1
1.0	0.17888889	0.5000000	0	1
1.0	0.20000000	0.5000000	0	1

ROC was used to select the optimal model using the largest value.
The final values used for the model were alpha = 0 and lambda = 0.2.

Nearest Shrunk Centroids:

```
> nsCTunedc
Nearest Shrunk Centroids

3333 samples
 14 predictor
  2 classes: 'yes', 'no'

Pre-processing: centered (14), scaled (14)
Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)
Summary of sample sizes: 2501, 2501, 2501, 2501, 2501, ...
Resampling results across tuning parameters:
```

threshold	ROC	Sens	Spec
0.0	0.5199949	0	1
0.1	0.5181676	0	1
0.2	0.5150843	0	1
0.3	0.5111367	0	1
0.4	0.5068965	0	1
0.5	0.5031414	0	1
0.6	0.5004422	0	1
0.7	0.4961718	0	1
0.8	0.4921142	0	1
0.9	0.4939239	0	1
1.0	0.4933488	0	1
1.1	0.4942849	0	1
1.2	0.4967039	0	1
1.3	0.4984841	0	1
1.4	0.5000000	0	1
1.5	0.5000000	0	1
1.6	0.5000000	0	1
1.7	0.5000000	0	1
1.8	0.5000000	0	1
1.9	0.5000000	0	1
2.0	0.5000000	0	1
2.1	0.5000000	0	1
2.2	0.5000000	0	1
2.3	0.5000000	0	1
2.4	0.5000000	0	1
2.5	0.5000000	0	1

```

2.6      0.5000000  0      1
2.7      0.5000000  0      1
2.8      0.5000000  0      1
2.9      0.5000000  0      1
3.0      0.5000000  0      1
3.1      0.5000000  0      1
3.2      0.5000000  0      1
3.3      0.5000000  0      1
3.4      0.5000000  0      1
3.5      0.5000000  0      1
3.6      0.5000000  0      1
3.7      0.5000000  0      1
3.8      0.5000000  0      1
3.9      0.5000000  0      1
4.0      0.5000000  0      1

```

ROC was used to select the optimal model using the largest value.
The final value used for the model was threshold = 0.

Model	ROC
Logistic Regression	0.5178
Linear Discriminant Analysis	0.5179
Partial Least Square Discriminant Analysis	0.5200
Penalized Model	0.5195
Nearest Shrunken Centroids	0.5199

The best model out of all of them is **Partial Least Square Discriminant Analysis (PLSDA)**.

R CODE:

```
install.packages(c("glmnet", "pamr", "rms", "sparseLDA", "subselect"))
```

```
#12.1
```

```
library(caret)
```

```
library(AppliedPredictiveModeling)
```

```
data(hepatic)
```

```
library(MASS)
```

```
set.seed(1)
```

```
barplot(table(injury), main="Imbalanced Class Distribution")
```

```
#PreProcess the data
```

```
#Removing nearzero variance
```

```
NV <- nearZeroVar(bio)
```

```
NV
```

```
noZVbio <- bio[,-NV]
```

```
noZVbio
```

```
#Calculating missing values
```

```
sum(is.na(noZVbio))
```

```
#Finding correlation
```

```
set.seed(1)
```

```
highCorBio<-findCorrelation(cor(noZVbio),cutoff = .80)
```

```
filteredCorBio <- noZVbio[,-highCorBio]
```

Split the data

set.seed(1)

trainingRows = createDataPartition(injury, p = .75, list= FALSE)

trainBio <- filteredCorBio[trainingRows,]

testBio <- filteredCorBio[-trainingRows,]

trainInjury <- injury[trainingRows]

testInjury <- injury[-trainingRows]

#Model building

##Multinomial Logistic Regression##

set.seed(1)

ctrl <- trainControl(summaryFunction = defaultSummary)

lrBio <- train(x=trainBio,

y = trainInjury,

method = "multinom",

preProc = c("center", "scale"),

metric = "Accuracy",

```
trControl = ctrl)
```

```
summary(lrBio)
```

```
plot(lrBio)
```

```
predictionLRBio<-predict(lrBio,testBio)
```

```
confusionMatrix(data =predictionLRBio,  
reference = testInjury)
```

```
##Linear Discriminant Analysis
```

```
set.seed(1)
```

```
ldaBio <- train(x = trainBio,  
y = trainInjury,  
method = "lda",  
preProc = c("center", "scale"),  
metric = "Accuracy",  
trControl = ctrl)
```



```
summary(ldaBio)
```

```
predictionLDABio <- predict(ldaBio,testBio)
```

```
confusionMatrix(data =predictionLDABio,  
                  reference = testInjury)
```

```
##### Partial Least Squares Discriminant Analysis #####
```

```
library(MASS)
```

```
set.seed(1)
```

```
plsFit <- train(x = trainBio,  
               y = trainInjury,  
               method = "pls",  
               tuneGrid = expand.grid(.ncomp = 1:1),  
               preProc = c("center","scale"),  
               metric = "Accuracy",  
               trControl = ctrl)
```

```
plsFit
```

```
plot(plsFit)
```

```
summary(plsFit)
```

```
varImp(plsFit, scale = FALSE)
```

```
predictionPLSBio <- predict(plsFit, testBio)
```

```
confusionMatrix(data = predictionPLSBio,  
                  reference = testInjury)
```

```
##### Penalized Models #####
```

```
## The primary package for penalized logistic regression is glmnet.
```

```
library(caret)
```

```
set.seed(1)
```

```
ctrl1 <- trainControl(method = "cv", number = 10)
```

```
glmnetGrid <- expand.grid(alpha = c(0, .1, .2, .4, .6, .8, 1),  
                          .lambda = seq(.01, .2, length = 10))
```

```
glmnetTuned <- train(x = trainBio, y = trainInjury, method = "glmnet",  
                    tuneGrid = glmnetGrid,  
                    preProc = c("center", "scale"),  
                    metric = "Accuracy", trControl = ctrl1)
```

```
summary(glmnetTuned)
```

```
important <- varImp(glmnetTuned, scale = FALSE)
```

```
plot(important, top = 5, scales = list(y = list(cex = .95)))
```

```
predictGlmnetBio <- predict(glmnTuned,testBio)
```

```
confusionMatrix(data = predictGlmnetBio,  
                 reference = testInjury)
```

```
plot(glmnTuned, plotType = "level")
```

```
##### Nearest Shrunken Centroids #####
```

```
library(pamr)
```

```
nscGrid <- data.frame(.threshold = seq(0,4, by=0.1))
```

```
set.seed(1)
```

```
nscTuned <- train(x = trainBio, y = trainInjury, method = "pam",
```

```
                 preProc = c("center", "scale"), tuneGrid = nscGrid,
```

```
                 metric = "Accuracy", trControl = ctrl)
```

```
nscTuned
```

```
plot(nscTuned)
```

```
summary(nscTuned)
```

```
predictNSC <-predict(nscTuned,testBio)
```

```
confusionMatrix(data =predictNSC,  
                  reference = testInjury)
```

```
predictors(nscTuned)
```

```
#12.3
```

```
#12.3
```

```
library(C50)
```

```
library(corrplot)
```

```
data(churn)
```

```
str(churnTrain)
```

```
table(churnTrain$churn)
```

```
#plot
```

```
par(mfrow = c(2,2))
```

```
plot(churn~state,data = churnTrain)

plot(churn~account_length,data = churnTrain)

plot(churn~area_code,data = churnTrain)

plot(churn~international_plan,data = churnTrain)

par(mfrow = c(2,2))

plot(churn~voice_mail_plan,data = churnTrain)

plot(churn~number_vmail_messages,data = churnTrain)

plot(churn~total_day_minutes,data = churnTrain)

plot(churn~total_day_calls,data = churnTrain)

par(mfrow = c(2,2))

plot(churn~total_day_charge,data = churnTrain)

plot(churn~total_eve_minutes,data = churnTrain)

plot(churn~total_eve_calls,data = churnTrain)

plot(churn~total_eve_charge,data = churnTrain)

par(mfrow = c(2,2))

plot(churn~total_night_minutes,data = churnTrain)

plot(churn~total_night_calls,data = churnTrain)

plot(churn~total_night_charge,data = churnTrain)

plot(churn~total_intl_minutes,data = churnTrain)

par(mfrow = c(1,3))

plot(churn~total_intl_calls,data = churnTrain)

plot(churn~total_intl_charge,data = churnTrain)

plot(churn~number_customer_service_calls,data = churnTrain)
```

#(c)

```
predict_train <- churnTrain[,-20]
```

```
ctrain <- churnTrain[,20]
```

```
predict_test <- churnTest[,-20]
```

```
ctest <- churnTest[,20]
```

#Dummy Variables

```
library(caret)
```

```
dummy <- dummyVars("~state + area_code + international_plan + voice_mail_plan",  
  data = predict_train, fullRank = TRUE)
```

```
dummytrain <- data.frame(predict(dummy, newdata = predict_train))
```

```
dummy <- dummyVars("~state + area_code + international_plan + voice_mail_plan",  
  data = predict_test, fullRank = TRUE)
```

```
dummytest <- data.frame(predict(dummy, newdata = predict_test))
```

Drop all factor predictors:

```
predict_train <- predict_train[,-c(1,3,4,5)]
```

```
predict_test <- predict_test[,-c(1,3,4,5)]
```

```
predict_train <- merge(predict_train, dummytrain, by =0)
```

```
predict_test <- merge(predict_test, dummytest, by =0)
```

```
predict_train <- predict_train[,-c(1)]
```

```
predict_test <- predict_test[,-c(1)]
```

```
#PreProcess
```

```
newdata <- preProcess(predict_train, method = c("center","scale","BoxCox"))
```

```
newdata_train <- predict(newdata, predict_train)
```

```
newdata_test <- predict(newdata, predict_test)
```

```
NV3 <- nearZeroVar(newdata_train)
```

```
newdata_train <- newdata_train[-NV3]
```

```
newdata_test <- newdata_test[-NV3]
```

```
highcor <- cor(newdata_train)
```

```
highcorpredict <- findCorrelation(highcor)
```

```
newdata_train <- newdata_train[,-highcorpredict]
```

```
newdata_test <- newdata_test[,-highcorpredict]
```

```
library(pROC)

ctrl1 = trainControl(method = "LGOCV",
                      summaryFunction=twoClassSummary,
                      classProbs=TRUE )
```

Logistic Regression Model:

```
set.seed(1)

lrnew <- train(x=newdata_train,
               y = ctrain,
               method = "glm",
               preProc = c("center", "scale"),
               metric = "ROC",
               trControl = ctrl1)
```

```
summary(lrnew)
```

lrnew

```
predictionLR<-predict(lrnew,newdata_test)
```



```
confusionMatrix(data =predictionLR,  
                 reference = ctest)
```

Linear Discriminant Analysis:

```
set.seed(1)
```

```
lda <- train(x = newdata_train,  
            y = ctrain,  
            method = "lda",  
            preProc = c("center", "scale"),  
            metric = "ROC",  
            trControl = ctrl1)
```

```
lda
```

```
summary(lda)
```

```
predictionLDA <- predict(lda,newdata_test)
```

```
confusionMatrix(data =predictionLDA,  
                 reference = ctest)
```

Partial Least Squares Discriminant Analysis

```
library(MASS)
```

```
set.seed(1)
```

```
plsFitc <- train(x = newdata_train,  
                y = ctrain,  
                method = "pls",  
                tuneGrid = expand.grid(.ncomp = 1:1),  
                preProc = c("center", "scale"),  
                metric = "ROC",  
                trControl = ctrl1)
```

```
plsFitc
```

```
summary(plsFitc)
```

```
varImp(plsFitc, scale = FALSE)
```

```
predictionPLS <- predict(plsFitc, newdata_test)
```

```
confusionMatrix(data = predictionPLS,  
                 reference = ctest)
```

Penalized Methods:

```
library(caret)
```

```
set.seed(1)
```

```
glmnetGrid <- expand.grid(.alpha = c(0, .1, .2, .4, .6, .8, 1),  
                          .lambda = seq(.01, .2, length = 10))
```

```
glmnetTuned <- train(x = newdata_train, y = ctrain, method = "glmnet",  
                    tuneGrid = glmnetGrid,  
                    preProc = c("center", "scale"),  
                    metric = "ROC", trControl = ctrl1)
```

```
glmnetTuned
```

```
important <- varImp(glmnetTuned, scale = FALSE)  
plot(important, top = 5, scales = list(y = list(cex = .95)))
```

```
predictGlmnet <- predict(glmnetTuned, newdata_test)  
confusionMatrix(data = predictGlmnet,  
                 reference = ctest)
```

```
plot(glmnetTuned, plotType = "level")
```

```
# Nearest shrunken Centroids:
```

```
library(pamr)
```

```
nscGrid <- data.frame(.threshold = seq(0,4, by=0.1))
```

```
set.seed(1)
```

```
nscTunedc <- train(x = newdata_train, y = ctrain, method = "pam",
```

```
  preProc = c("center", "scale"), tuneGrid = nscGrid,
```

```
  metric = "ROC", trControl = ctrl1)
```

```
nscTunedc
```

```
plot(nscTunedc)
```

```
summary(nscTunedc)
```

```
predictNSCc <- predict(nscTunedc, newdata_test)
```

```
confusionMatrix(data = predictNSCc,
```

```
  reference = ctest)
```

```
predictors(nscTunedc)
```